

## Running Batch Jobs on UIUC's Campus Cluster

**Access** (Via Unix/Linux Shell)

(In Windows, use **PuTTY** or **PowerShell SSH-Sessions**.)

```
ssh -Y -l user cc-login.campuscluster.illinois.edu      # log in to head nodes
scp ./local user@ccftp1.campuscluster.illinois.edu:~/remote  # copy data to CC
scp user@ccftp1.campuscluster.illinois.edu:~/remote ./local  # copy data from CC
```

### Module Commands

module	command interface	load <i>module</i>	load <i>module</i> (set \$PATH etc.)
avail	list available modules	unload <i>module</i>	unload <i>module</i>
list	list loaded modules	swap <i>mod1 mod2</i>	swap two modules
show <i>module</i>	show commands w/i <i>module</i>	purge	unload all modules

### Batch Commands (Moab/Torque)

qstat	query status of batch jobs	qdel <i>jobid</i>	delete job
-u \$USER	limit to all jobs of \$USER	qpeek <i>jobid</i>	view stdout and stderr of job
-n <i>jobid</i>	list nodes alloc. to job	qhold <i>jobid</i>	delay a job in queue
-f <i>jobid</i>	show details of job	qrls <i>jobid</i>	release a job to queue
-t <i>jobid</i> []	give status of all jobs in array	showstart <i>jobid</i>	show current estimated start time of job
js <i>jobid</i>	shows batch job script		

### Job Submission Flags (qsub)

#### BASIC JOB CONTROL

-l walltime= <i>hh:mm:ss</i>	specify desired job run time	-q <i>queue</i>	specify <i>queue</i>
-l nodes= <i>n</i> :ppn=12	request <i>n</i> nodes and 12 processors per node	-N <i>jobname</i>	name job
-d <i>dir</i>	open job shell in directory <i>dir</i>	-I	interactive shell

#### OUTPUT & NOTIFICATION

-j oe	pipe stdout and stderr to	-m abe	send mail at <u>b</u> eginning, <u>e</u> nd,
-o <i>a.log</i>	<i>a.log</i>	-M <i>a@b.com</i>	and on job <u>a</u> bort to <i>a@b.com</i>

#### ADVANCED JOB CONTROL

-t <i>a-b</i>	specify job array	-l nodes= <i>n</i> :ppn=12 :m96G	request memory (not CSE queue)
-W depend=?	introduce dependency (run before or after another job)	-v	export environment variables to script
		-x	enable X-forwarding

## EFFECTIVE QUEUE UTILIZATION

Requested times should be close to expected run times for scheduling efficiency. There is also a limit to how many simultaneous jobs you can run—the others will sit in the queue until you have more available slots.

*With fairshare scheduling, one can adopt different strategies to run jobs quickly on Campus Cluster:*

- Submit smaller jobs which will backfill between larger jobs as nodes become available.
- Submit either job arrays or fewer larger jobs with independent processes within them (using &).
- Utilize the **test** and **debug** queues when appropriate.
- Jobs shorter than four hours will currently run on either Taub or Golub as resources become available.

*Responsible queue usage involves requesting only the nodes you need and using them responsibly.*

- Monitor your jobs to make sure that they don't hang (crash without terminating the process) (**qpeek**).
- Submit jobs as close to the necessary size as possible—this makes the queueing algorithm more efficient.
- Only use interactive jobs to edit, compile, and build your programs as necessary—don't let them sit idle.

## Moab Environment Variables

<code>\$PBS_JOBID</code>	Batch queue job ID number
<code>\$PBS_NODEFILE</code>	List of nodes available for batch job
<code>\$PBS_O_WORKDIR</code>	Directory from which batch job was submitted

## Executing an MPI Program (Internode, Distributed Memory)

The number of processes available to MPI generally equals the requested number of nodes × cores per node.

```
module load mvapich2
```

```
mpiexec /path/to/application ./inputfile &> ./outputfile.$PBS_JOBID
```

## Executing an OpenMP Program (Intranode, Shared Memory)

```
export OMP_NUM_THREADS=12
```

```
./hello_world_openmp
```

## Example PBS Script

```
#!/bin/csh
#PBS -l walltime=00:05:00
#PBS -l nodes=2:ppn=12
#PBS -q test
#PBS -N hello_world_mpi
#PBS -j oe
#PBS -o hello_world_mpi.log.$PBS_JOBID
#PBS -m abe
#PBS -M $USER@illinois.edu

setenv job "hello_world_mpi"
setenv RUN_DIR $HOME/src/mpi-mwe/cpp/
module load mvapich2/1.6-intel

# Document settings for debugging purposes (also can use ldd, echo $PATH).
set echo ; module list ; which mpiexec
mpiexec $RUN_DIR/$job
```

## Automating Access (alias, ssh)

It can become tedious to type in the domain every time you wish to connect. Add this to your ~/.bash\_profile:

```
alias taub='ssh -Y -l username cc-login.campuscluster.illinois.edu'
```

You can also set up your machine to authenticate with a private key rather than inputting a password every time.

```
cd ; mkdir .ssh ; ssh-keygen -t rsa # use any passphrase (but don't leave it empty)
```

```
chmod +600 .ssh/id_rsa # make your private key only readable by you
```

```
scp .ssh/id_rsa.pub <username>@<remote>: # include the ':'
```

```
ssh -l username cc-login.campuscluster.illinois.edu
```

```
mkdir .ssh # this is on the remote machine
```

```
cat id_rsa.pub >> .ssh/authorized_keys # append ('>>') rather than replace
```

```
chmod 700 .ssh # make your public key only accessible by you
```

Now log out and test this by sshing back in again. Both of these need be done only once (on each machine).